

Xiang Zhou
Ph.D. Student
<https://owenzx.github.io>

April 20, 2021
Mobile: +1-(984) 215-7399
Email: xzh@cs.unc.edu

EDUCATION

AUG. 2018–PRESENT	Ph.D. student in Computer Science University of North Carolina at Chapel Hill, US
SEPT. 2014–JULY. 2018	B.S.E. in Computer Science, IEEE Honor Class Shanghai Jiao Tong University, China

EXPERIENCE

AUG. 2018–PRESENT	UNC-NLP Research Group, UNC Chapel Hill <i>Research Assistant, Supervised by Prof. Mohit Bansal</i> <ul style="list-style-type: none">• Unsupervised Part-Of-Speech Tagging• Analysis of instability and uncertainty in Natural Language Inference/Question Answering models• Robustifying Natural Language Inference models with adversarial methods
JUNE 2020–OCT. 2020	Amazon Alexa AI <i>Applied Scientist Intern, Supervised by Heba Elfardy and Tom Butler</i> <ul style="list-style-type: none">• Automatic detection of unreliable news
JULY 2016–JUNE 2018	Speech Lab, Shanghai Jiao Tong University <i>Research Assistant, Supervised by Prof. Kai Yu</i> <ul style="list-style-type: none">• RL-based framework to bootstrap neural dialogue policy from human/statistical models• Joint optimization of dialogue policy and dialogue state tracker
JULY 2017–SEPT. 2017	NLP Group, University of Notre Dame <i>Research Intern, Supervised by Prof. David Chiang</i> <ul style="list-style-type: none">• Character-level neural language models

PUBLICATIONS

-
- | | |
|------|--|
| 2021 | Xiang Zhou* , Yixin Nie* and Mohit Bansal. <i>Distributed NLI: Learning to Predict Human Opinion Distributions on Language Reasoning</i> . Preprint on arXiv |
| 2021 | Xiang Zhou , Heba Elfardy, Christos Christodoulopoulos, Thomas Butler and Mohit Bansal. <i>Hidden Biases in Unreliable News Detection Datasets</i> . In Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics (EACL 2021) |
| 2020 | Xiang Zhou , Yixin Nie, Hao Tan and Mohit Bansal. <i>The Curse of Performance Instability in Analysis Datasets: Consequences, Source, and Suggestions</i> . In Proceedings of the Conference on Empirical Methods in Natural Language Processing 2020 (EMNLP 2020) |
| 2020 | Yixin Nie, Xiang Zhou and Mohit Bansal. <i>What Can We Learn from Collective Human Opinions on Natural Language Inference Data?</i> . In Proceedings of the Conference on Empirical Methods in Natural Language Processing 2020 (EMNLP 2020) |
| 2020 | Xiang Zhou and Mohit Bansal. <i>Towards Robustifying NLI Models Against Lexical Dataset Biases</i> . In Proceedings of the 2020 Annual Conference of the Association for Computational Linguistics (ACL 2020) |

- 2017 Lu Chen, **Xiang Zhou**, Cheng Chang, Runzhe Yang and Kai Yu. *Rule-Guided Safe and Efficient On-line Dialogue Policy Learning*. In Proceedings of the Conference on Empirical Methods in Natural Language Processing 2017 (EMNLP 2017)
- 2017 Cheng Chang, Runzhe Yang, Lu Chen, **Xiang Zhou** and Kai Yu. *Affordable On-line Dialogue Policy Learning*. In Proceedings of the Conference on Empirical Methods in Natural Language Processing 2017 (EMNLP 2017)
- 2017 Lu Chen, Runzhe Yang, Cheng Chang, Zihao Ye, **Xiang Zhou** and Kai Yu. *On-line Dialogue Policy Learning with Companion Teaching*. In Proceedings of the 15th Conference of the European Chapter of the Association for Computational Linguistics (EACL 2017)

PROFESSIONAL SERVICES

CONFERENCE REVIEWER EMNLP 2020, EACL 2021, NAACL 2021, ACL 2021
 WORKSHOP REVIEWER HAMLETS@NeurIPS 2020, RobustML@ICLR 2021

ACADEMIC ACTIVITIES

SEPT. 2017–JULY. 2018 Co-translator of the Chinese translation of *Reinforcement Learning: An Introduction*

SELECTED SCHOLARSHIPS AND AWARDS

2016 Eleme Scholarship (Top 10%)
 2015, 2016 Academic Excellence Scholarship Prize B (Top 10%)
 2015 Xindong Scholarship (Top 10%)
 2015 The Mathematical Contest in Modeling 2015 Meritorious Winner

SKILLS

Programming Languages: Python, C++, Lua, MATLAB, \LaTeX , Verilog HDL
 Machine Learning Frameworks: PyTorch, TensorFlow, MXNet, DyNet