

**Xiang Zhou**  
Ph.D. Student  
<https://owenzx.github.io>

April 16, 2022  
Mobile: +1-(984) 215-7399  
Email: [xzh@cs.unc.edu](mailto:xzh@cs.unc.edu)

## EDUCATION

---

AUG. 2018–PRESENT	Ph.D. student in Computer Science <b>University of North Carolina at Chapel Hill, US</b>
SEPT. 2014–JULY. 2018	B.S.E. in Computer Science, IEEE Honor Class <b>Shanghai Jiao Tong University, China</b>

## EXPERIENCE

---

AUG. 2018–PRESENT	<b>UNC-NLP Research Group, UNC Chapel Hill</b> <i>Research Assistant, Supervised by Prof. Mohit Bansal</i> <ul style="list-style-type: none"><li>Improving compositional generalization and robustness of NLP models</li><li>Unsupervised Part-Of-Speech Tagging</li><li>Analysis of instability and uncertainty in Natural Language Inference models</li></ul>
MAY 2021–DEC. 2021	<b>Facebook AI Research</b> <i>Research Intern, Supervised by Jean Maillard</i> <ul style="list-style-type: none"><li>Incorporating syntax in NMT</li></ul>
JUNE 2020–OCT. 2020	<b>Amazon Alexa AI</b> <i>Applied Scientist Intern, Supervised by Heba Elfardy and Tom Butler</i> <ul style="list-style-type: none"><li>Automatic detection of unreliable news</li></ul>
JULY 2016–JUNE 2018	<b>Speech Lab, Shanghai Jiao Tong University</b> <i>Research Assistant, Supervised by Prof. Kai Yu</i> <ul style="list-style-type: none"><li>RL-based framework to bootstrap neural dialogue policy from human/statistical models</li><li>Joint optimization of dialogue policy and dialogue state tracker</li></ul>
JULY 2017–SEPT. 2017	<b>NLP Group, University of Notre Dame</b> <i>Research Intern, Supervised by Prof. David Chiang</i> <ul style="list-style-type: none"><li>Character-level neural language models</li></ul>

## PUBLICATIONS

- 
- |      |   |
|------|---|
| 2022 | <b>Xiang Zhou</b> , Shiyue Zhang and Mohit Bansal. <i>Masked Part-Of-Speech Model: Does modeling long context help unsupervised POS-tagging?</i> In Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL 2022)                     |
| 2022 | <b>Xiang Zhou*</b> , Yixin Nie* and Mohit Bansal. <i>Distributed NLI: Learning to Predict Human Opinion Distributions on Language Reasoning</i> . In Findings of the Association for Computational Linguistics: ACL 2022 (ACL 2022 Findings)  |
| 2022 | Archiki Prasad, Peter Hase, <b>Xiang Zhou</b> and Mohit Bansal. <i>GRIPS: Gradient-free, Edit-based Instruction Search for Prompting Large Language Models</i> . In Arxiv   |
| 2021 | <b>Xiang Zhou</b> , Heba Elfardy, Christos Christodoulopoulos, Thomas Butler and Mohit Bansal. <i>Hidden Biases in Unreliable News Detection Datasets</i> . In Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics (EACL 2021) ( <b>Best Long Paper Honorable Mention</b> ) |

- 2020 **Xiang Zhou**, Yixin Nie, Hao Tan and Mohit Bansal. *The Curse of Performance Instability in Analysis Datasets: Consequences, Source, and Suggestions*. In Proceedings of the Conference on Empirical Methods in Natural Language Processing 2020 (EMNLP 2020)
- 2020 Yixin Nie, **Xiang Zhou** and Mohit Bansal. *What Can We Learn from Collective Human Opinions on Natural Language Inference Data?*. In Proceedings of the Conference on Empirical Methods in Natural Language Processing 2020 (EMNLP 2020)
- 2020 **Xiang Zhou** and Mohit Bansal. *Towards Robustifying NLI Models Against Lexical Dataset Biases*. In Proceedings of the 2020 Annual Conference of the Association for Computational Linguistics (ACL 2020)
- 2017 Lu Chen, **Xiang Zhou**, Cheng Chang, Runzhe Yang and Kai Yu. *Rule-Guided Safe and Efficient On-line Dialogue Policy Learning*. In Proceedings of the Conference on Empirical Methods in Natural Language Processing 2017 (EMNLP 2017)
- 2017 Cheng Chang, Runzhe Yang, Lu Chen, **Xiang Zhou** and Kai Yu. *Affordable On-line Dialogue Policy Learning*. In Proceedings of the Conference on Empirical Methods in Natural Language Processing 2017 (EMNLP 2017)
- 2017 Lu Chen, Runzhe Yang, Cheng Chang, Zihao Ye, **Xiang Zhou** and Kai Yu. *On-line Dialogue Policy Learning with Companion Teaching*. In Proceedings of the 15th Conference of the European Chapter of the Association for Computational Linguistics (EACL 2017)

## PROFESSIONAL SERVICES

---

REVIEWER	ACL Rolling Review
CONFERENCE REVIEWER	EMNLP 2020, EACL 2021, NAACL 2021, ACL 2021, AAAI 2022
JOURNAL REVIEWER	Computational Linguistics (CL)
WORKSHOP REVIEWER	HAMLETS@NeurIPS 2020, RobustML@ICLR 2021, Insights@EMNLP 2021

## ACADEMIC ACTIVITIES

---

SEPT. 2017–JULY. 2018 Co-translator of the Chinese translation of *Reinforcement Learning: An Introduction*

## SELECTED SCHOLARSHIPS AND AWARDS

---

2021	Best Long Paper Award Honorable Mention, EACL 2021
2016	Eleme Scholarship (Top 10%)
2015, 2016	Academic Excellence Scholarship Prize B (Top 10%)
2015	Xindong Scholarship (Top 10%)

## SKILLS

---

Programming Languages:	Python, C++, Lua, MATLAB, $\text{\LaTeX}$ , Verilog HDL
Machine Learning Frameworks:	PyTorch, TensorFlow, MXNet, DyNet