

Xiang Zhou
Ph.D. Student
<https://owenzx.github.io>

November 20, 2022
Mobile: +1-(984) 215-7399
Email: xzh@cs.unc.edu

EDUCATION

- | | |
|----------------------------------|---|
| AUG. 2018–MAY 2023
(EXPECTED) | Ph.D. student in Computer Science
University of North Carolina at Chapel Hill, US |
| SEPT. 2014–JULY. 2018 | B.S.E. in Computer Science, IEEE Honor Class
Shanghai Jiao Tong University, China |

EXPERIENCE

- | | |
|----------------------|--|
| AUG. 2018–PRESENT | UNC-NLP Research Group, UNC Chapel Hill
<i>Research Assistant, Supervised by Prof. Mohit Bansal</i> <ul style="list-style-type: none">Improving compositional generalization and robustness of NLP modelsUnsupervised Part-Of-Speech TaggingAnalysis of instability and uncertainty in Natural Language Inference models |
| MAY 2022–AUG. 2022 | Google Assistant
<i>Research Intern, Supervised by Manaal Faruqui, Aditya Gupta and Shyam Upadhyay</i> <ul style="list-style-type: none">Sequential instruction understanding with limited supervisions. |
| MAY 2021–DEC. 2021 | Facebook AI Research
<i>Research Intern, Supervised by Jean Maillard</i> <ul style="list-style-type: none">Incorporating syntax in NMT. |
| JUNE 2020–OCT. 2020 | Amazon Alexa AI
<i>Applied Scientist Intern, Supervised by Heba Elfardy and Tom Butler</i> <ul style="list-style-type: none">Automatic detection of unreliable news. |
| JULY 2016–JUNE 2018 | Speech Lab, Shanghai Jiao Tong University
<i>Research Assistant, Supervised by Prof. Kai Yu</i> <ul style="list-style-type: none">RL-based framework to bootstrap neural dialogue policy from human/statistical models.Joint optimization of dialogue policy and dialogue state tracker. |
| JULY 2017–SEPT. 2017 | NLP Group, University of Notre Dame
<i>Research Intern, Supervised by Prof. David Chiang</i> <ul style="list-style-type: none">Character-level neural language models. |

PUBLICATIONS

- | | |
|------|---|
| 2022 | Yichen Jiang, Xiang Zhou* , and Mohit Bansal. <i>Mutual exclusivity training and primitive augmentation to induce compositionality</i> . In Proceedings of the Conference on Empirical Methods in Natural Language Processing 2022 (EMNLP 2022) |
| 2022 | Xiang Zhou , Shiyue Zhang and Mohit Bansal. <i>Masked Part-Of-Speech Model: Does modeling long context help unsupervised POS-tagging?</i> In Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL 2022) |
| 2022 | Xiang Zhou* , Yixin Nie* and Mohit Bansal. <i>Distributed NLI: Learning to Predict Human Opinion Distributions on Language Reasoning</i> . In Findings of the Association for Computational Linguistics: ACL 2022 (ACL 2022 Findings) |

- 2022 Archiki Prasad, Peter Hase, **Xiang Zhou** and Mohit Bansal. *GRIPS: Gradient-free, Edit-based Instruction Search for Prompting Large Language Models*. Preprint on arXiv
- 2021 **Xiang Zhou**, Heba Elfardy, Christos Christodoulopoulos, Thomas Butler and Mohit Bansal. *Hidden Biases in Unreliable News Detection Datasets*. In Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics (EACL 2021) (**Best Long Paper Honorable Mention**)
- 2020 **Xiang Zhou**, Yixin Nie, Hao Tan and Mohit Bansal. *The Curse of Performance Instability in Analysis Datasets: Consequences, Source, and Suggestions*. In Proceedings of the Conference on Empirical Methods in Natural Language Processing 2020 (EMNLP 2020)
- 2020 Yixin Nie, **Xiang Zhou** and Mohit Bansal. *What Can We Learn from Collective Human Opinions on Natural Language Inference Data?*. In Proceedings of the Conference on Empirical Methods in Natural Language Processing 2020 (EMNLP 2020)
- 2020 **Xiang Zhou** and Mohit Bansal. *Towards Robustifying NLI Models Against Lexical Dataset Biases*. In Proceedings of the 2020 Annual Conference of the Association for Computational Linguistics (ACL 2020)
- 2017 Lu Chen, **Xiang Zhou**, Cheng Chang, Runzhe Yang and Kai Yu. *Rule-Guided Safe and Efficient On-line Dialogue Policy Learning*. In Proceedings of the Conference on Empirical Methods in Natural Language Processing 2017 (EMNLP 2017)
- 2017 Cheng Chang, Runzhe Yang, Lu Chen, **Xiang Zhou** and Kai Yu. *Affordable On-line Dialogue Policy Learning*. In Proceedings of the Conference on Empirical Methods in Natural Language Processing 2017 (EMNLP 2017)
- 2017 Lu Chen, Runzhe Yang, Cheng Chang, Zihao Ye, **Xiang Zhou** and Kai Yu. *On-line Dialogue Policy Learning with Companion Teaching*. In Proceedings of the 15th Conference of the European Chapter of the Association for Computational Linguistics (EACL 2017)

PROFESSIONAL SERVICES

REVIEWER	ACL Rolling Review
CONFERENCE REVIEWER	EMNLP (2020, 2022), EACL 2021, NAACL 2021, ACL 2021, AAAI (2022, 2023)
JOURNAL REVIEWER	Computational Linguistics (CL)
WORKSHOP REVIEWER	HAMLETS@NeurIPS 2020, RobustML@ICLR 2021, Insights@EMNLP 2021, MMNLU@EMNLP 2022

VOLUNTARY ACTIVITIES

SEPT. 2017–JULY. 2018 Participated in the Chinese translation of *Reinforcement Learning: An Introduction*

SELECTED SCHOLARSHIPS AND AWARDS

2021	Best Long Paper Award Honorable Mention, EACL 2021
2016	Eleme Scholarship (Top 10%)
2015, 2016	Academic Excellence Scholarship Prize B (Top 10%)
2015	Xindong Scholarship (Top 10%)

SKILLS

Machine Learning Frameworks:	PyTorch, JAX+Flax, TensorFlow
Programming Languages:	Python, C/C++